

INTRODUCTION TO MACHINE LEARNING

K-NEAREST NEIGHBOR ALGORITHM


Mingon Kang, PhD
Department of Computer Science@UNLV

KNN


- K-Nearest Neighbors (KNN)
- Simple, but a very powerful classification algorithm
- Classifies based on a **similarity measure**
- Non-parametric
- Lazy learning
 - ▣ Does not “learn” until the test example is given
 - ▣ Whenever we have a new data to classify, we find its K-nearest neighbors from the training data

KNN: Classification Approach

- Classified by “**MAJORITY VOTES**” for its neighbor classes
- ▣ Assigned to the most common class amongst its K -nearest neighbors (by measuring “distance” between data)




KNN: Example



KNN: Pseudocode

- Step 1: Determine parameter K = number of nearest neighbors
- Step 2: Calculate the distance between the query-instance and all the training examples.
- Step 3: Sort the distance and determine nearest neighbors based on the k -th minimum distance.
- Step 4: Gather the category Y of the nearest neighbors.
- Step 5: Use simple majority of the category of nearest neighbors as the prediction value of the query instance.

KNN: Example




KNN: Euclidean distance matrix

Table 1. Euclidean distance matrix D listing all possible pairwise Euclidean distances between 19 samples.

	x ₁	x ₂	x ₃	x ₄	x ₅	x ₆	x ₇	x ₈	x ₉	x ₁₀	x ₁₁	x ₁₂	x ₁₃	x ₁₄	x ₁₅	x ₁₆	x ₁₇	x ₁₈	x ₁₉
x ₂	1.5																		
x ₃	1.4	1.6																	
x ₄	1.6	1.4	1.3																
x ₅	1.7	1.4	1.5	1.5															
x ₆	1.3	1.4	1.4	1.5	1.4														
x ₇	1.6	1.3	1.4	1.4	1.5	1.8													
x ₈	1.5	1.4	1.6	1.3	1.7	1.6	1.4												
x ₉	1.4	1.3	1.4	1.5	1.2	1.4	1.3	1.5											
x ₁₀	2.3	2.4	2.5	2.3	2.6	2.7	2.8	2.7	3.1										
x ₁₁	2.9	2.8	2.9	3.0	2.9	3.1	2.9	3.1	3.0	1.5									
x ₁₂	3.2	3.3	3.2	3.1	3.3	3.4	3.3	3.4	3.5	3.3	1.6								
x ₁₃	3.3	3.4	3.2	3.2	3.3	3.4	3.2	3.3	3.5	3.6	1.4	1.7							
x ₁₄	3.4	3.2	3.5	3.4	3.7	3.5	3.6	3.3	3.5	3.6	1.5	1.8	0.5						
x ₁₅	4.2	4.1	4.1	4.1	4.1	4.1	4.1	4.1	4.1	4.1	1.7	1.6	0.3	0.5					
x ₁₆	4.1	4.1	4.1	4.1	4.1	4.1	4.1	4.1	4.1	4.1	1.6	1.5	0.4	0.5	0.4				
x ₁₇	5.9	6.2	6.2	5.8	6.1	6.0	6.1	5.9	5.8	6.0	2.3	2.3	2.5	2.3	2.4	2.5			
x ₁₈	6.1	6.3	6.2	5.8	6.1	6.0	6.1	5.9	5.8	6.0	3.1	2.7	2.6	2.3	2.5	2.6	3.0		
x ₁₉	6.0	6.1	6.2	5.8	6.1	6.0	6.1	5.9	5.8	6.0	3.0	2.9	2.7	2.4	2.5	2.8	3.1	0.4	


Decision Boundaries

- Voronoi diagram
 - ▣ Describes the areas that are nearest to any given point, given a set of data.
 - ▣ Each line segment is equidistant between two points of opposite class



Decision Boundaries


- With large number of examples and possible noise in the labels, the decision boundary can become nasty!
- ▣ “Overfitting” problem




Effect of K

- Larger k produces smoother boundary effect
- When $K=N$, always predict the majority class

$K=1$



$K=15$




Figures from Hastie, Tibshirani and Friedman (Elements of Statistical Learning)

Discussion

- Which model is better between $K=1$ and $K=15$?
- Why?

How to choose k?

□ Empirically optimal k?



Pros and Cons

□ Pros

- ▣ Learning and implementation is extremely simple and Intuitive
- ▣ Flexible decision boundaries

□ Cons

- ▣ Irrelevant or correlated features have high impact and must be eliminated
- ▣ Typically difficult to handle high dimensionality
- ▣ Computational costs: memory and classification time computation

Similarity and Dissimilarity

□ Similarity

- Numerical measure of how alike two data objects are.
- Is higher when objects are more alike.
- Often falls in the range $[0,1]$

□ Dissimilarity

- Numerical measure of how different are two data objects
- Lower when objects are more alike
- Minimum dissimilarity is often 0
- Upper limit varies

□ Proximity refers to a similarity or dissimilarity

Euclidean Distance


□ Euclidean Distance

$$dist = \sqrt{\sum_{k=1}^p (a_k - b_k)^2}$$

Where p is the number of dimensions (attributes) and a_k and b_k are, respectively, the k -th attributes (components) or data objects a and b .

□ Standardization is necessary, if scales differ.

Euclidean Distance



point	x	y
p1	0	2
p2	2	0
p3	3	1
p4	5	1

	p1	p2	p3	p4
p1	0	2.828	3.162	5.099
p2	2.828	0	1.414	3.162
p3	3.162	1.414	0	2
p4	5.099	3.162	2	0

Minkowski Distance

- Minkowski Distance is a generalization of Euclidean Distance

$$dist = \sum_{k=1}^p |a_k - b_k|^r$$

Where r is a parameter, p is the number of dimensions (attributes) and a_k and b_k are, respectively, the k -th attributes (components) or data objects a and b

Minkowski Distance: Examples

- $r = 1$. City block (Manhattan, taxicab, L_1 norm) distance.
 - ▣ A common example of this is the Hamming distance, which is just the number of bits that are different between two binary vectors
- $r = 2$. Euclidean distance
- $r \rightarrow \infty$. “supremum” (L_{max} norm, L_∞ norm) distance.
 - ▣ This is the maximum difference between any component of the vectors
- Do not confuse r with p , i.e., all these distances are defined for all numbers of dimensions.

Cosine Similarity

- If d_1 and d_2 are two document vectors

$$\cos(d_1, d_2) = (d_1 \cdot d_2) / (||d_1|| ||d_2||),$$

Where \cdot indicates vector dot product and $||d||$ is the length of vector d .

- Example:

$$d_1 = 3 \ 2 \ 0 \ 5 \ 0 \ 0 \ 0 \ 2 \ 0 \ 0$$

$$d_2 = 1 \ 0 \ 0 \ 0 \ 0 \ 0 \ 0 \ 1 \ 0 \ 2$$

$$d_1 \cdot d_2 = 3*1 + 2*0 + 0*0 + 5*0 + 0*0 + 0*0 + 0*0 + 2*1 + 0*0 + 0*2 = 5$$

$$||d_1|| = (3*3 + 2*2 + 0*0 + 5*5 + 0*0 + 0*0 + 0*0 + 2*2 + 0*0 + 0*0)^{0.5} = (42)^{0.5} = 6.481$$

$$||d_2|| = (1*1 + 0*0 + 0*0 + 0*0 + 0*0 + 0*0 + 0*0 + 1*1 + 0*0 + 2*2)^{0.5} = (6)^{0.5} = 2.245$$

$$\cos(d_1, d_2) = .3150$$

Cosine Similarity

□ $\cos(d_1, d_2) = \begin{cases} 1: \text{exactly the same} \\ 0: \text{orthogonal} \\ -1: \text{exactly opposite} \end{cases}$

Feature scaling

- Standardize the range of independent variables (features of data)
- A.k.a Normalization or Standardization

Standardization

- Standardization or Z-score normalization
 - ▣ Rescale the data so that the mean is zero and the standard deviation from the mean (standard scores) is one

$$X_{norm} = \frac{X - \mu}{\sigma}$$

μ is mean, σ is a standard deviation from the mean
(standard score)

Min-Max scaling

- Scale the data to a fixed range – between 0 and 1

$$X_{\text{morm}} = \frac{X - X_{\text{min}}}{X_{\text{max}} - X_{\text{min}}}$$

Efficient implementation

- Consider data as a matrix or a vector
- Matrix/Vector computational is much more efficient than computing with loop

Discussion

- Can we use KNN for regression problems?